

# CS63 Spring 2018

## Classifying Recipes By Cuisines: Random Forests and Neural Networks

Muha Haque and Emilie Shepherd

May 14, 2018

### 1 Introduction

The goal of this experiment was to correctly classify the cuisine of a recipe, given a list of the recipe's ingredients. We approached this classification problem in two different ways, with a Random Forest ensemble learner, and with a Neural Network.

Random forest classifiers are an ensemble learning method in which a designated number of decision trees are fitted to samples of the data set and then averaged together. They use bagging, in which the data set is sampled with replacement. The decision trees then classify each of the samples from the data sets, and a plurality vote is taken over the predictions that are made by the different classifiers. Random forests are an appropriate machine learning algorithm for our experiment because decision trees could be prone to overfit to our complex data set, but the bagging and plurality voting averages out the different biases and variances of the decision trees and controls overfitting as a result.

Artificial Neural Networks are a deep learning method designed to mirror the connectivity and learning ability of the human brain. A neural network is a directed, weighted graph with distinct layers of nodes. The network “learns” by passing information through the layers, comparing the network's predicted output to the correct output, and then adjusting the edge weights by propagating the error back through the network. The neural network we used in this experiment had two types of layers: Densely connected and Dropout. Every node in a dense layer is connected to every node in the following layer. Especially in larger nets, this can cause the net to overfit to the training data. To combat this, we used dropout layers, which selectively silence nodes to prevent overfitting. Neural networks are adept at finding patterns in data with a large number of features, which made them an appropriate choice for this experiment.

### 2 Method and Details

#### 2.1 Data Set

The data used was provided by Yummly.com for a Kaggle competition. The data set was a .json file of recipes, each of which consisted of a list of ingredients labeled with one of 20 cuisines. The data set contained 39,774 recipes, which we divided into a training set ( $n=31,774$ ) and a test set ( $n=8,000$ ). The number of recipes per cuisine varied widely, from 7,838 (Italian) to 467 (Russian)

Cuisine	Train Set	Test Set
Brazilian	372	95
British	655	149
Cajun Creole	1211	335
Chinese	2142	531
Filipino	607	148
French	2125	521
Greek	933	242
Indian	2409	594
Irish	546	121
Italian	6276	1562
Jamaican	417	109
Japanese	1148	275
Korean	641	189
Mexican	5154	1284
Moroccan	674	147
Russian	397	92
Southern US	3402	918
Spanish	780	209
Thai	1225	314
Vietnamese	660	165

Table 1: Recipe breakdown by cuisine.

(Table 1). There were 6206 unique ingredients in the training data set, and 508 ingredients that appeared only in the test set.

## 2.2 Preprocessing

### 2.2.1 Manual Cleaning

Because the ingredients list was pulled directly from user input on Yummly, it had a fair amount of inconsistencies and redundancies – for example, “light mayonnaise,” “reduced fat mayonnaise,” and “Hellmann’s Light Mayonnaise” all appeared separately on the ingredients list. We preprocessed the data by overwriting some ingredient variants with a more general token word. We selected a few common ingredients and created lists of equivalent variations (Table 2), and then replaced all instances of those variations in the data set with their respective token words. The most successful version of cleaned data replaced 135 ingredients with 21 token words.

We trained on both the original data and the preprocessed data. The first version of cleaned data produced a lower accuracy than the original data, but after editing the lists of replaced ingredients we were able to produce an increase in accuracy.

### 2.2.2 One-Hot Encoding

Because our data was discrete and non-numerical, we needed to transform it into a form more easily understood by the Random Forest and Neural Net. Before training, we performed a one-hot

Flour Tortillas
“flour tortillas”
“Old El Paso Flour Tortillas”
“large flour tortillas”
“fajita size flour tortillas”
“flour tortillas (not low fat)”
“soft taco size flour tortillas”
“low-fat flour tortillas”
“Azteca flour tortillas”

Table 2: An example entry in the dictionary used to clean the data.

encoding on the input data. We extracted a list of all ingredients in the training set and transformed each recipe into a binary encoding of that list, with a 1 at the index of each ingredient in the recipe, and a 0 at all other indices. This binary vector was simple to input to both the Random Forest Classifier and the Neural Network.

## 2.3 Implementation Details

### 2.3.1 Random Forests: scikit-learn

We used scikit-learn for our implementation of random forests. Specifically, we used the RandomForestClassifier in the Ensemble library. The most important parameters that we were able to set and test were the number of decision trees in the forest built before voting and averaging and the number of features to consider when the trees look for splits. The tradeoff we had to consider with the number of trees was the general increase in accuracy with more trees but the extra time it took for the code to run as a result. With the number of features, runtime was a huge consideration because each unique ingredient was a feature. Additionally, considering too large a number of features would also risk reducing the diversity of individual decision trees. Our model produced the highest accuracy with 500 trees in the forest.

### 2.3.2 Neural Networks: Keras

For our neural networks, we used the Keras implementation. Keras allowed us to test our networks quickly and also to easily build our network with layers on top of another using the Sequential model. The only layers we needed were their Dense Layers, which were just basic densely connected layers where we could set the number of units and the activation function applied to the units in a layer, as well as the input shape for the input layer; and their Dropout layers, where we could select the percentage of units from the prior layer to set to 0. We used the reLu activations for nearly all of our layers except for our output layer, for which we used the softmax activation function which is often ideal for classification tasks.

In our experiments, we also chose to vary the optimizer used by the network and the loss function. The optimizers we tried were stochastic gradient descent, Adam, and Adagrad; Adam and Adagrad are both gradient descent optimization algorithms. The loss function we mainly used was categorical cross-entropy, but we also had one experiment in which we used mean squared error. We set accuracy as our metric for all of the networks that we tested.

### 2.3.3 Training methods

Both models used the same input format for training and testing but required the outputs to be formatted slightly differently. The random forest fit the one-hot encoded training input to the training output of cuisine labels, and then predicted the cuisine labels for our testing inputs, which were also one-hot encoded. For the neural network, we had to take the extra step of using scikit-learn's Label Encoder to turn our cuisine labels into integer categorical outputs that the neural network we used required. Each network we used for our experiment trained for 10 epochs using the one-hot encoding of the training recipes and the label-encoded training outputs, and was then evaluated using the one-hot encoding of the testing recipes and their label-encoded outputs.

## 3 Results

### 3.1 Random Forest Model

Our random forest implementation showed the most significant increases in its predictive accuracy when we increased the number of trees, as shown in Figure ?? . The highest accuracy rate we achieved was 72.71%, with 500 trees in our forest. This result was achieved with our modified data set.




Figure 1: Graph showing the accuracy of the random forest model with different numbers of trees.

Our random forest performed worse on the original data set, achieving an accuracy of 71.78%. None of the other parameters we changed helped the forest's predictive ability significantly, and changing them often led to worse results. The parameter that was closest to being relevant was




Figure 2: Graph showing testing accuracy and training accuracy for the neural network. The large difference between the two lines indicates extreme overfitting.

the max features parameter; we ultimately used sci-kit learn’s default, which used the square root of the total features, to achieve our best result. The closest we were able to get while changing the parameter was when we considered 10% of the features, which increased our accuracy slightly with smaller sets of trees, but decreased performance when the number of trees was greater than 100; with 500 trees, it reached an accuracy of only 70.6%. Our best results with the random forest model put us at a ranking of about 1028/1388 in the Kaggle competition.

### 3.2 Neural Network Model

For the neural network, the highest predictive accuracy we achieved with the cleaned data set was 78.35%. This neural network had 5 dense layers, with 1024, 512, 256, 128, and 20 units from the input later to the output layer. All these layers used the ReLU activation function, with the exception of the output layer which used the softmax activation function. Between each dense layer was a dropout layer with a dropout rate of 0.4. The loss function was categorical cross-entropy and the optimizer was Adagrad; Adam had slightly worse performance with a peak accuracy of 77.6%, while stochastic gradient descent was slightly worse still with a peak accuracy of 75.35%. We trained our network for 50 epochs. Notably, while the predictive accuracy on the test set increased slowly and eventually leveled out during the 50 epochs, our neural network started to overfit extremely very quickly, eventually reaching accuracies of 99% on the training data during epochs, as shown in Figure ???. We increased the dropout rates as much as we could, and found that after a rate of 0.4 the accuracy on the test set dropped; for example, the same network with

dropout rates of 0.5 had an accuracy of 77.85% and a minimal decrease in overfitting. With the original data, the same neural net achieved an accuracy of 78.16%. Our best neural network’s results put us at a ranking of about 641/1388 in the Kaggle competition.

## 4 Conclusions

### 4.1 Discussion of Findings

Given a large set of recipes and the cuisine categories they belonged to, we aimed to build 2 models for predicting the cuisine types of a test set of recipes. By using one-hot encoding, we were able to vectorize our complex, non-numerical input set and make it usable for both the random forest and the neural network. We cleaned the original data set to reduce the number of ingredients in recipes which were essentially the same but were counted as different features, and both of our models generally predicted better as a result. Our data preprocessing produced a larger increase in accuracy in the random forest model than in the neural network model.

The best accuracy of our random forest model was 72.71%, with cleaned data and 500 trees in the forest. This level of accuracy would likely not have been possible with a single decision tree, because the large number of features in our data set would be difficult to differentiate without extreme overfitting.

The neural network performed notably better than the random forest, with a best accuracy of 78.35%. However, this was at the price of heavy overfitting; this is likely due to the large number of features, the 6000+ ingredients. Changing the optimizer from stochastic gradient descent to Adagrad notably increased the accuracy, potentially because Adagrad maintains different learning rates for each of its parameters unlike SGD and better captures the complexity and wide-range of our features as a result.

### 4.2 Looking to the Future

We believe that a higher accuracy could be achieved with better preprocessing of the data. Our “cleaned” data set still contained a large number of redundant ingredients. Additionally, we could investigate using scikit-learn’s DictVectorizer or some other feature extractor pre-training. Most of our data processing was done by hand, and there are almost certainly more sophisticated methods available. For the neural network, we could look into other ways to reduce the overfitting of our model without decreasing our predictive accuracy in addition to adding dropout layers; given the large set of input features, this might also lie in finding some other way to succinctly vectorize our data without losing the complex differences between recipes of different cuisines.